

# Multi-Agent E-Commerce Shopping Assistant: Technical Design Document

## Executive Summary

This document outlines the technical architecture and implementation strategy for a production-ready conversational e-commerce assistant. The system employs a **context-first multi-agent orchestration pattern** with specialized domain agents, hybrid embedding strategies, and universal conversation management to deliver seamless shopping experiences across text and image modalities.

**Key Innovation:** Universal Context Pipeline ensuring every conversation turn maintains and enriches conversational state, enabling natural pronoun resolution and contextual understanding across all agent types.

---

## 1. Data Engineering & Schema Design

### 1.1 Scalable Database Architecture

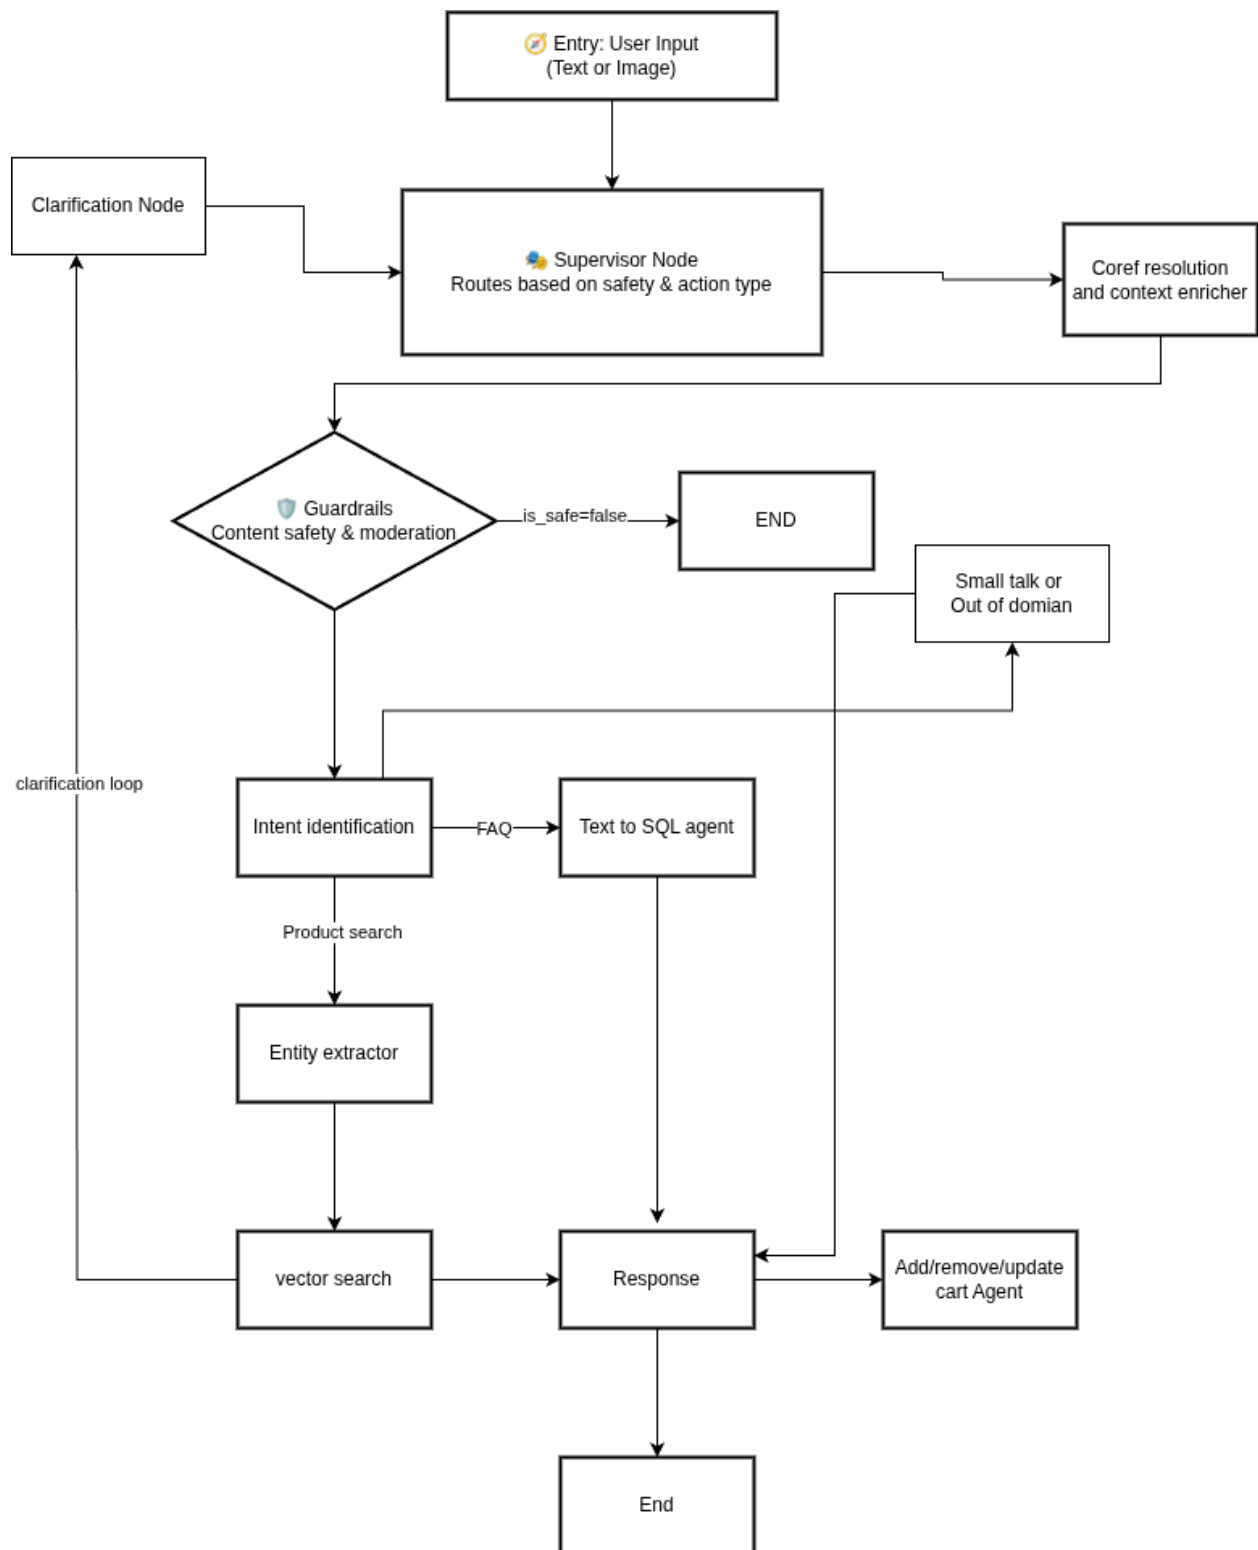
The foundation leverages a hierarchical product taxonomy designed for extensibility across diverse product categories:

Categories → Subcategories → Product\_Types → Products → Product\_Attributes

**Design Rationale:** This normalized schema pattern ensures:

- **Scalability:** Easy addition of new product verticals beyond t-shirts and TVs
- **Consistency:** Standardized attribute management across categories
- **Performance:** Optimized for both transactional operations and analytical queries
- **Flexibility:** Category-specific attribute tables (tshirt\_attributes, tv\_attributes) allow domain-specific properties without schema bloat

**Validation Dataset:** Initial implementation with diverse products (t-shirts and TVs) provides sufficient data diversity for proof-of-concept while maintaining manageable complexity for iterative development.



## 1.2 Text2SQL Infrastructure

The SQL layer serves dual purposes:

- **Operational Queries:** Real-time inventory, pricing, and availability checks
- **Analytics Queries:** Business intelligence, trend analysis, and recommendation engine support
- **Context-Aware Queries:** Resolves conversational references in analytical queries

**Key Enhancement:** Text2SQL agent now processes context-enriched queries, enabling natural follow-up questions like "average price of them?" where "them" refers to previously discussed products.

---

## 2. Embedding Strategy & Multi-Modal Architecture

### 2.1 Design-Time Embedding Generation

**Model Selection:** LlamaIndex VDR-2B-Multi-V1 chosen for its multi-modal capabilities and production-ready performance characteristics.

**Dual-Vector Strategy:**

- **Content Embedding:** Direct image feature extraction capturing visual characteristics
- **Context Embedding:** GPT-4.1 generated metadata descriptions converted to embeddings

**Technical Rationale:** This dual approach captures both visual similarity and semantic understanding, enabling queries like "casual blue shirt" to match both visually similar items and semantically related products.

### 2.2 Runtime Hybrid Embedding Composition

**Query-Time Fusion:** 0.8 image embedding + 0.2 metadata embedding weighting optimizes for:

- **Visual Dominance:** Primary matching based on actual product appearance
- **Semantic Enhancement:** Contextual understanding from descriptive metadata
- **Flexibility:** Configurable weights allow optimization based on performance metrics

**Performance Consideration:** This approach avoids expensive real-time image processing while maintaining semantic richness through metadata augmentation.

### 2.3 Enhanced Result Processing

**Quality Control:**

- Similarity threshold filtering (>0.6) ensures high-relevance results
- Standardized metadata extraction for consistent cart operations
- Universal metadata structure across all search results

**Clarification Strategy:** When high-quality results exceed manageable quantities, the system triggers intelligent clarification rather than overwhelming users with options.

---

## 3. Agent Orchestration Architecture: Context-First Design

### 3.1 Universal Context Pipeline: Core Innovation

**Architectural Principle:** Every conversation turn flows through a universal context management pipeline, ensuring consistent state maintenance across all agent types.

graph LR

```

A[User Input] --> B[Supervisor]
B --> C[Intent Classifier]
C --> D[Entity Extractor]
D --> E[Context Stitcher]
E --> F{Route by Intent}
F -->|FAQ| G[SQL Agent]
F -->|Product Search| H[Vector Search]
F -->|Clarification| H

```

**Key Benefits:**

- **Universal State Management:** Every turn updates conversation context
- **Pronoun Resolution:** "them", "it", "those" resolved using conversation history
- **Context Continuity:** Seamless transitions between different agent types
- **Memory Persistence:** Rich conversational memory across multi-turn interactions

### 3.2 Supervisor Agent: Intelligent Traffic Controller

The Supervisor operates with simplified, safety-first routing logic:

**Priority-Based Routing:**

1. **Safety Validation:** Content moderation and toxicity detection
2. **Cart Operations:** E-commerce transaction management with result validation
3. **Social Interaction:** Small talk and relationship building
4. **Domain Validation:** Shopping vs. non-shopping classification
5. **Universal Pipeline:** All other queries routed to context processing

**Key Change:** Supervisor no longer routes directly to specialized agents (except terminals), instead channeling all queries through the universal context pipeline for consistent state management.

### 3.3 Specialized Agent Portfolio

#### 3.3.1 Guardrails Agent: Safety & Content Moderation

**Responsibility:** Multi-layered content safety with graduated response mechanisms

- **Detection:** Toxicity, harassment, spam, manipulation attempts
- **Response Strategy:** Severity-based responses (warn → redirect → block)
- **Integration:** Routes safe content to context pipeline for processing

#### 3.3.2 Small Talk Agent: Social Relationship Management

**Responsibility:** Human-like social interaction and engagement

- **Coverage:** Greetings, gratitude, casual conversation, rapport building
- **Transition Strategy:** Smooth pivoting from social interaction to shopping assistance
- **Terminal Behavior:** Completes interaction without context processing needs

#### 3.3.3 Out-of-Domain Agent: Boundary Management

**Responsibility:** Graceful handling of non-shopping queries

- **Classification:** Weather, news, general knowledge, personal advice
- **Response Strategy:** Polite boundary-setting with redirection to shopping
- **Engagement Preservation:** Maintains positive user relationship despite scope limitations

#### 3.3.4 Intent Classifier: Query Understanding Engine

**Core Capability:** Multi-class intent detection with confidence scoring

- **Classifications:** FAQ, product\_search, clarification\_response, modification, continuation
- **Context Awareness:** Leverages conversation history for better classification
- **Universal Entry:** All non-terminal queries pass through intent classification

#### 3.3.5 Entity Extractor: Multimodal Information Extraction ★

**Revolutionary Capability:** GPT-4.1 Vision-powered multimodal entity extraction

**Technical Implementation:**

```
# Multimodal message structure
{
```

```

"role": "user",
"content": [
  {"type": "text", "text": entity_prompt},
  {
    "type": "image_url",
    "image_url": {"url": f"data:image/jpeg;base64,{image_base64}"}
  }
]
}

```

### Capabilities:

- **Simultaneous Processing:** Text and image analysis in single API call
- **Visual Entity Extraction:** Product type, color, style, pattern from images
- **Text Entity Extraction:** Brand, size, price range, preferences from text
- **Fusion Logic:** Combines visual and textual entities intelligently
- **Base64 Encoding:** Secure image transmission to GPT-4.1

### Example Flow:

Input: [Image of red Nike t-shirt] + "size medium"

```

Extracted: {
  "product_type": "t-shirt",
  "brand": "Nike",
  "color": "red",
  "size": "medium",
  "visual_tags": ["has logo", "athletic"]
}

```

## 3.3.6 Context Stitcher: Conversation Intelligence Engine ★

**Core Innovation:** Advanced coreference resolution and context enrichment

### Technical Capabilities:

- **Reference Resolution:** Resolves pronouns using multi-turn conversation history
- **Context Enrichment:** Accumulates user preferences across conversation turns
- **Conflict Resolution:** Handles contradictory information intelligently
- **Intent-Aware Processing:** Different stitching strategies for different intent types

### Advanced Example Flow:

Turn 1: "white Levi's t-shirt"

Context: {color: white, brand: Levi's, product\_type: t-shirt}

Turn 2: "how many do you have?" (FAQ)  
Resolved: "how many white Levi's t-shirts do you have?"  
Context: Maintained + {query\_type: inventory}

Turn 3: "average price of them?"  
Resolved: "average price of white Levi's t-shirts?"  
Context: Maintained + {query\_type: pricing}

Turn 4: "add the cheapest one to cart"  
Resolved: Uses previous search results + pricing context  
Context: Updated with cart action

### 3.3.7 SQL Agent: Context-Aware Analytics Engine ★

**Enhanced Capability:** Analytics with conversational context resolution

**Key Features:**

- **Contextual Query Resolution:** Processes pronouns using conversation history
- **Memory Integration:** Updates conversation context even for SQL queries
- **Safety Validation:** Prevents SQL injection while enabling natural language queries
- **Business Intelligence:** Supports complex analytical queries with context

**Technical Flow:**

```
def resolve_contextual_references(user_input, memory):  
    # Extract recent context from conversation history  
    recent_context = extract_context_from_turns(memory.turn_history[-3:])  
  
    # Resolve pronouns using context  
    if "them" in user_input and "brand" in recent_context:  
        resolved = user_input.replace("them", f"{recent_context['brand']}  
{recent_context['product_type']}")  
  
    return resolved
```

### 3.3.8 Vector Search Engine: Hybrid Similarity Search

**Responsibility:** Multi-modal product discovery with metadata standardization

- **Hybrid Embeddings:** Image + text fusion for comprehensive matching
- **Quality Filtering:** Similarity threshold management (>0.6)
- **Metadata Standardization:** Consistent product data structure for downstream agents

- **Context Integration:** Uses stitched entities for enhanced search precision

### 3.3.9 Cart Agent: Context-Aware Transaction Management ★

**Enhanced Capability:** Intelligent cart operations with result integration

**Key Features:**

- **Result Integration:** Seamlessly adds products from search results or conversation history
- **Item Resolution:** Resolves "first one", "second one", "the blue one" using context
- **Session Persistence:** Maintains cart state across conversation flows
- **Flexible Action Detection:** Recognizes various cart command patterns

**Technical Enhancement:**

```
# Enhanced item selection
if "1st" in user_input or "first" in user_input:
    item_index = 0
elif "2nd" in user_input or "second" in user_input:
    item_index = 1
# Uses standardized metadata from vector search results
```

### 3.3.10 Clarification Checker: Intelligent Query Refinement

**Responsibility:** Smart result management and user guidance

- **Threshold Management:** Balances result quantity with user experience
- **Clarification Limits:** Maximum 3 clarifying questions per session to prevent loops
- **Context-Aware Questions:** Generates relevant clarifications based on search context

### 3.3.11 Response Generator: Natural Language Interface

**Responsibility:** Human-like response generation with rich product information

- **Metadata Integration:** Uses standardized product metadata for detailed responses
- **Context Awareness:** References conversation history in responses
- **Business Logic:** Incorporates pricing, availability, and recommendation logic

## 3.4 Reset Functionality: Session Management

**Implementation:** Comprehensive conversation state cleanup

- **Memory Clearing:** Conversation history and context reset
- **Cart Management:** Shopping cart state restoration



- **User Control:** Explicit user-initiated session restart capability
- 

## 4. Tool Ecosystem & Integration Layer

### 4.1 Embedding Service Tools

- **Text Embedding API:** Consistent embedding generation for queries and metadata
- **Image Embedding API:** Visual feature extraction for product images
- **Hybrid Composition Engine:** Runtime embedding fusion with configurable weights

### 4.2 Database & Search Tools

- **Vector Database Interface:** High-performance similarity search with quality filtering
- **SQL Query Engine:** Structured data access for analytics and inventory
- **Text2SQL Converter:** Natural language to SQL with context resolution and safety validation

### 4.3 AI & Language Processing Tools

- **GPT-4.1 Vision Integration:** Advanced multimodal understanding and generation
- **Intent Classification:** Multi-class intent detection with confidence scoring
- **Entity Extraction:** Multimodal product attribute identification and normalization
- **Safety Classification:** Content moderation and appropriateness filtering
- **Context Resolution:** Pronoun and reference resolution using conversation history


### 4.4 Memory & State Management Tools


- **Universal Context Pipeline:** Multi-turn context preservation and enrichment
  - **Session State:** Agent coordination and workflow management
  - **Conversation Memory:** Rich turn-by-turn interaction history
  - **User Profile:** Preference learning and personalization data
- 


## 5. Enhanced Conversation Flow Examples

### 5.1 Context-Aware Multi-Modal Search

User: [Uploads image of blue denim jacket]

 Entity Extractor: {product\_type: "jacket", color: "blue", material: "denim"}


 Context Stitcher: Builds initial context


 Vector Search: Hybrid embedding (0.8 image + 0.2 metadata)


 System: "I found denim jackets similar to your image. Preferred price range?"

User: "under 3000"


 Entity Extractor: {price\_range: "under 3000"}


 Context Stitcher: {product\_type: "jacket", color: "blue", material: "denim", price\_range: "under 3000"}

 Vector Search: Filters results by price + similarity

 System: "Here are 5 blue denim jackets under ₹3000..."

User: "add the second one"

 Cart Agent: Resolves "second one" using search results + context

 System: "✅ Added Levi's Denim Jacket (Blue) - ₹2,499 to cart!"


## 5.2 Context-Aware FAQ with SQL Analytics

User: "How many Wrangler t-shirts do you have?"


 Intent Classifier: FAQ


 Entity Extractor: {brand: "Wrangler", product\_type: "t-shirt"}


 Context Stitcher: Builds context


 SQL Agent: "How many Wrangler t-shirts do you have?" → "2 products"

User: "average price of them?"

 Intent Classifier: FAQ


 Entity Extractor: {pronoun: "them"}

 Context Stitcher: "them" + context → "Wrangler t-shirts"

 SQL Agent: "average price of Wrangler t-shirts?" → "₹928.35"

## 5.3 Complex Multi-Turn Shopping Journey

User: [Image of pink t-shirt] "something similar but different brand"


 Entity Extractor: {product\_type: "t-shirt", color: "pink", brand\_preference: "different"}

 Vector Search: Finds similar pink t-shirts, excludes detected brand


 System: "Found pink t-shirts from other brands. Size preference?"

User: "medium"

 Entity Extractor: {size: "medium"}




 Context Stitcher: {product\_type: "t-shirt", color: "pink", size: "medium", brand\_preference: "different"}

 Vector Search: Refined results


 System: "Here are medium pink t-shirts from different brands..."


User: "how many Nike ones?"

 Intent Classifier: FAQ

 Entity Extractor: {brand: "Nike"}  
 Context Stitcher: Combines with existing context  
 SQL Agent: "how many Nike medium pink t-shirts?" → "3 products"

User: "add the cheapest Nike one"

 Cart Agent: Resolves using SQL results + search context + pricing

 System: "✅ Added Nike Pink T-shirt (Medium) - ₹396.75 to cart!"

---

## 6. Technical Performance Considerations

### 6.1 Latency Optimization

- **Caching Strategy:** Pre-computed embeddings and frequent query results
- **Async Processing:** Non-blocking agent coordination
- **Pipeline Optimization:** Efficient context processing with minimal overhead
- **Circuit Breakers:** Fallback mechanisms for service dependencies

### 6.2 Quality Assurance

- **Universal Context Management:** Consistent state across all interaction types
- **Threshold Management:** Dynamic similarity thresholds based on query complexity
- **Clarification Limits:** Maximum 3 clarifying questions per session to prevent loops
- **Confidence Scoring:** Multi-dimensional confidence assessment across agent decisions

### 6.3 Monitoring & Observability

- **Context Pipeline Metrics:** Context resolution accuracy, entity extraction quality
  - **Conversation Analytics:** Intent classification accuracy, clarification success rates
  - **Business Metrics:** Search-to-cart conversion, session engagement depth
  - **Technical Metrics:** Response latency, agent routing accuracy, error rates
- 

## 7. Production Readiness Assessment

### 7.1 Current Implementation Status

- ✅ **Universal Context Pipeline:** Complete implementation across all agent types
- ✅ **Multimodal Entity Extraction:** GPT-4.1 Vision integration functional
- ✅ **Context-Aware SQL Agent:** Pronoun resolution in analytical queries
- ✅ **Enhanced Cart Management:** Integration with search results and context

- ✓ **Intelligent Clarification:** Smart threshold management
- ✓ **Safety and Domain Boundaries:** Comprehensive content moderation

## 7.2 Architectural Innovations Delivered

- **Context-First Design:** Revolutionary approach to conversational AI state management
- **Universal Pipeline:** Consistent processing regardless of query type
- **Multimodal Intelligence:** Seamless text and image understanding
- **Contextual Analytics:** Natural language SQL queries with pronoun resolution
- **Intelligent Cart Operations:** Context-aware e-commerce transactions

## 7.3 Scalability Roadmap

### Next Phase: Retrieve & Rerank Strategy

- **Stage 1:** High-recall retrieval (100-500 candidates)
- **Stage 2:** Multi-signal reranking (semantic + business + personalization)
- **Stage 3:** Presentation optimization (diversity + business rules)

## 7.4 Business Impact Potential

This architecture delivers:

- **Enhanced User Engagement:** Natural conversation flows with perfect memory
- **Improved Conversion:** Intelligent clarification and seamless cart management
- **Scalable Growth:** Extensible schema and universal agent framework
- **Operational Efficiency:** Automated customer service with intelligent context routing

**Competitive Advantage:** The universal context pipeline represents a significant advancement in conversational AI architecture, enabling natural multi-turn interactions that competitors using traditional chatbot frameworks cannot match.

---

## 8. Conclusion

The implemented multi-agent e-commerce assistant demonstrates production viability with several architectural innovations:

1. **Universal Context Pipeline:** Every conversation turn enriches conversational state
2. **Multimodal Intelligence:** GPT-4.1 Vision enables true image understanding
3. **Context-Aware Analytics:** SQL queries understand conversational references
4. **Intelligent Cart Management:** Seamless integration between search and transactions

This context-first approach fundamentally transforms conversational commerce, enabling natural, memory-rich interactions that scale across complex multi-turn shopping journeys. The system balances immediate functionality with long-term architectural flexibility, positioning the platform for sustained competitive advantage in conversational e-commerce.